



Flight Price Prediction Project Report



Submitted by:

Anurag Raj

ACKNOWLEDGMENT

I have utilized a few external resources that helped me to complete the project. I ensured that I learn from the samples and modify things according to my project requirement. All the external resources that were used in creating this project are listed below:

- 1) <https://www.google.com/>
- 2) <https://www.youtube.com/>
- 3) https://scikit-learn.org/stable/user_guide.html
- 4) <https://github.com/>
- 5) <https://www.kaggle.com/>
- 6) <https://medium.com/>
- 7) <https://towardsdatascience.com/>
- 8) <https://www.analyticsvidhya.com/>

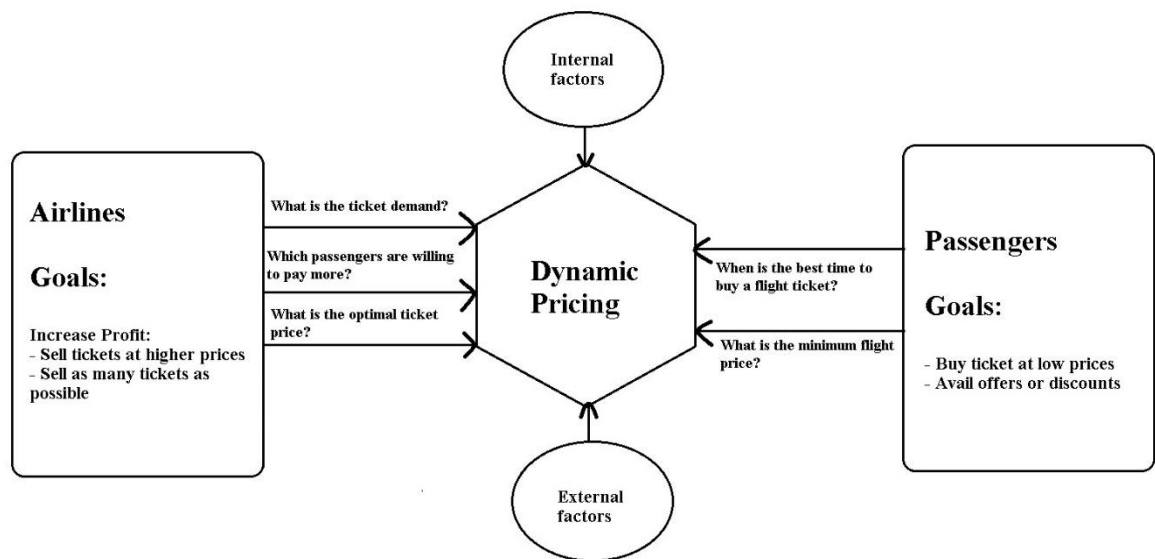
INTRODUCTION

- Business Problem Framing

The airline industry is considered as one of the most sophisticated industry in using complex pricing strategies. Nowadays, ticket prices can vary dynamically and significantly for the same flight, even for nearby seats. The ticket price of a specific flight can change up to 7 times a day. Customers are seeking to get the lowest price for their ticket, while airline companies are trying to keep their overall revenue as high as possible and maximize their profit. However, mismatches between available seats and passenger demand usually leads to either the customer paying more or the airlines company losing revenue. Airlines companies are generally equipped with advanced tools and capabilities that enable them to control the pricing process. However, customers are also becoming more strategic with the development of various online tools to compare prices across various airline companies. In addition, competition between airlines makes the task of determining optimal pricing is hard for everyone.

Anyone who has booked a flight ticket knows how unexpectedly the prices vary. The cheapest available ticket on a given flight gets more and less expensive over time. This usually happens as an attempt to maximize revenue based on

1. Time of purchase patterns (making sure last-minute purchases are expensive)
2. Keeping the flight as full as they want it (raising prices on a flight which is filling up in order to reduce sales and hold back inventory for those expensive last-minute expensive purchases)



- **Conceptual Background of the Domain Problem**

Airline companies use complex algorithms to calculate flight prices given various conditions present at that particular time. These methods take financial, marketing, and various social factors into account to predict flight prices. Nowadays, the number of people using flights has increased significantly. It is difficult for airlines to maintain prices since prices change dynamically due to different conditions. That's why we will try to use machine learning to solve this problem. This can help airlines by predicting what prices they can maintain. It can also help customers to predict future flight prices and plan their journey accordingly.

- **Review of Literature**

As per the requirement of client, I have scrapped the data from online sites and based on that data I have did analysis like for based on which feature of my data prices are changing. and checked the relationship of flight price with all the feature like what flight he should choose.

- **Motivation for the Problem Undertaken**

I have worked on this on the bases of client requirements and followed all the steps till model deployment.

Analytical Problem Framing

- Mathematical/ Analytical Modeling of the Problem

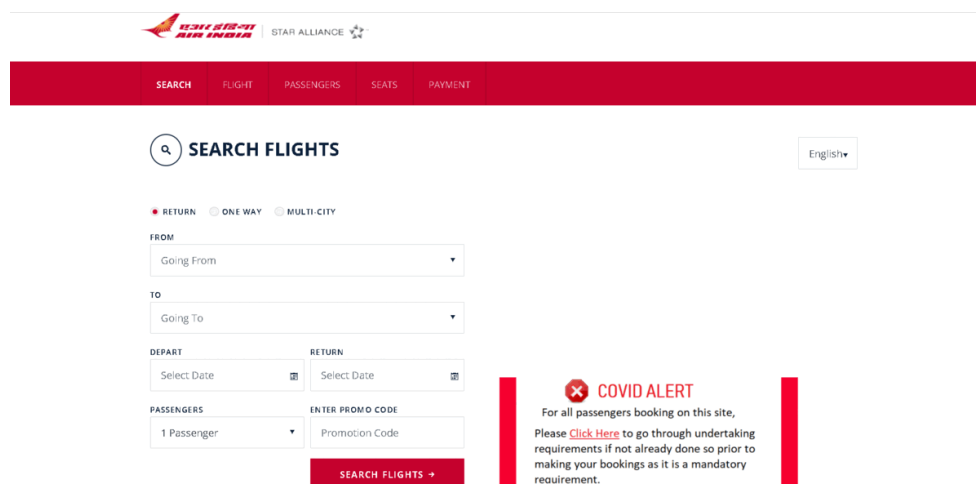
In our scrapped dataset, our target variable "Flight_Prices " is a continuous variable. Therefore, we will be handling this modelling problem as regression.

This project is done in three parts:

- Data Collection
- Data Analysis
- Model Building

1. Data Collection

You have to scrape at least 1500 rows of data. You can scrape more data as well, it's up to you, More the data better the model. In this section you have to scrape the data of flights from different websites (yatra.com, skyscanner.com, official websites of airlines, etc). The number of columns for data doesn't have limit, it's up to you and your creativity. Generally, these columns are airline name, date of journey, source, destination, route, departure time, arrival time, duration, total stops and the target variable price. You can make changes to it, you can add or you can remove some columns, it completely depends on the website from which you are fetching the data.



The screenshot shows the Air India website's flight search interface. At the top, there's a red navigation bar with the Air India logo and 'STAR ALLIANCE' text. Below this is a red bar with navigation links: SEARCH, FLIGHT, PASSENGERS, SEATS, and PAYMENT. The main search area is titled 'SEARCH FLIGHTS' with a magnifying glass icon. It includes radio buttons for 'RETURN' (selected), 'ONE WAY', and 'MULTI-CITY'. There are input fields for 'FROM' (Going From) and 'TO' (Going To). Below these are 'DEPART' and 'RETURN' date selectors. There's a 'PASSENGERS' dropdown set to '1 Passenger' and an 'ENTER PROMO CODE' field. A red 'SEARCH FLIGHTS +>' button is at the bottom. A 'COVID ALERT' banner is visible on the right side of the search area.

2. Data Analysis

After cleaning the data, you have to do some analysis on the data. Do airfares change frequently? Do they move in small increments or in large jumps? Do they tend to go up or down over time? What is the best time to buy so that the consumer can save the most by taking the least risk? Does price increase as we get near to departure date? Is Indigo cheaper than Jet Airways? Are morning flights expensive?

3. Model Building

After collecting the data, you need to build a machine learning model. Before model building do all data pre-processing steps. Try different models with different hyper parameters and select the best model.

Follow the complete life cycle of data science. Include all the steps like

1. Data Cleaning
2. Exploratory Data Analysis
3. Data Pre-processing
4. Model Building
5. Model Evaluation
6. Selecting the best model

- Data Sources and their formats

The dataset is in the form of CSV (Comma Separated Value) format and consists of 9 columns (8 features and 1 label) with 5805 number of records as explained below:

- Airline_Names : This shows the list of all the Airline Names for which the data got scraped
- Departure_Time : In this column we have the timings of every flight departure

- **Arrival_Time** : Here in this column we have the timings of every flight arrival
- **Flight_Duration** : We can see the total duration of a flight that it took to fly from the source to the destination
- **Source_Place** : Gives us the name of the source place where the flight journey began
- **Destination_Place** : Shows us the name of the destination place where the flight journey ended
- **Meal_Availability** : Provides us with the information on type of meal that the passenger is eligible for
- **Number_Of_Stops** : Lists the number of stops the flight is going to take to complete the entire journey
- **Flight_Prices** : Finally we have our label column that has the ticket prices for the aircraft journey

We can see our dataset includes a target label "Used Car Price" column and the remaining feature columns can be used to determine or help in predicting the price of the used cars. Since price is a continuous value it makes this to be a Regression problem!

```
In [2]: df = pd.read_csv("C:\\Users\\anura\\OneDrive\\Desktop\\anurag\\project\\Flight_Price_Data.csv")
```

I am importing the collected dataset comma separated values file and storing it into our dataframe for further usage.

```
In [3]: df # checking the first 5 and last 5 rows
```

```
Out[3]:
```

	Airline_Names	Departure_Time	Arrival_Time	Flight_Duration	Source_Place	Destination_Place	Meal_Availability	Number_Of_Stops	Flight_Prices
0	Air Asia	12:40	20:15	7h 35m	New Delhi	Mumbai	No Meal Fare	1 Stop	5,953
1	Air Asia	11:55	20:15	8h 20m	New Delhi	Mumbai	No Meal Fare	1 Stop	5,953
2	Air Asia	16:15	06:20	14h 05m	New Delhi	Mumbai	No Meal Fare	1 Stop	5,953
3	Go First	18:50	20:45	1h 55m	New Delhi	Mumbai	No Meal Fare	Non Stop	5,954
4	Go First	09:05	11:05	2h 00m	New Delhi	Mumbai	No Meal Fare	Non Stop	5,954
...
5800	Air India	08:55	08:20	23h 25m	Lucknow	Jaipur	No Meal Fare	1 Stop	9,302
5801	Air India	08:55	09:20	24h 25m	Lucknow	Jaipur	No Meal Fare	2 Stop(s)	16,287
5802	Air India	14:45	09:20	18h 35m	Lucknow	Jaipur	No Meal Fare	2 Stop(s)	16,865
5803	Air India	08:55	09:20	24h 25m	Lucknow	Jaipur	No Meal Fare	2 Stop(s)	16,865
5804	Air India	15:30	09:20	17h 50m	Lucknow	Jaipur	Free Meal	3 Stop(s)	19,749

5805 rows x 9 columns


Here we are taking a look at the first 5 and last 5 rows of our dataset. It shows that we have a total of 5805 rows and 9 columns present in our dataframe. In the above cell we can see our dataset includes a target label "Flight_Prices" column and the remaining feature columns can be used to determine or help in predicting the price of the flights for different places. Since price is a continuous value it makes this to be a Regression problem!

- Data Preprocessing Done

For the data pre-processing step, I checked through the dataframe for missing values and renamed values that needed a better meaningful name.

```
df.isna().sum() # checking for missing values
```

```
Airline_Names      0
Departure_Time     0
Arrival_Time       0
Flight_Duration    0
Source_Place       0
Destination_Place  0
Meal_Availability  0
Number_Of_Stops    0
Flight_Prices      0
dtype: int64
```



Checked the datatype details for each column to understand the numeric ones and its further conversion process.


```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5805 entries, 0 to 5804
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline_Names          5805 non-null   object
1   Departure_Time         5805 non-null   object
2   Arrival_Time           5805 non-null   object
3   Flight_Duration        5805 non-null   object
4   Source_Place           5805 non-null   object
5   Destination_Place      5805 non-null   object
6   Meal_Availability       5805 non-null   object
7   Number_Of_Stops        5805 non-null   object
8   Flight_Prices          5805 non-null   object
dtypes: object(9)
memory usage: 408.3+ KB
```

I also took a look at all the unique value present in each of the columns and then decided to view the details for those column values that had less than 10 categories.

```
df.nunique().sort_values().to_frame("Unique Values")
```

Unique Values	
Meal_Availability	3
Number_Of_Stops	5
Airline_Names	6
Source_Place	9
Destination_Place	9
Departure_Time	223
Arrival_Time	232
Flight_Duration	409
Flight_Prices	1571

The various data processing performed on our data set are shown below with the code.

```
# Meal_Availability
df.Meal_Availability.replace({"No Meal Fare": "No Meals", "Free Meal": "Free Meals", "eCash 250": "eCash Meals"},
                             inplace = True)
df["Meal_Availability"].value_counts()
```

```
# Number_Of_Stops
```

```
df.Number_Of_Stops.replace({"Non Stop": 0, "1 Stop": 1, "2 Stop(s)": 2, "3 Stop(s)": 3, "4 Stop(s)": 4},  
                           inplace = True)  
df["Number_Of_Stops"].value_counts()
```

```
# Departure_Time
```

```
df["Dep_Hour"] = pd.to_datetime(df.Departure_Time, format="%H:%M").dt.hour  
df["Dep_Min"] = pd.to_datetime(df.Departure_Time, format="%H:%M").dt.minute  
df["Departure_Time"] = df['Dep_Hour'] + df['Dep_Min'] / 60  
#df.drop(columns = ['Dep_Hour', 'Dep_Min'], inplace=True)  
df.head()
```

```
# Arrival_Time
```

```
df["Arr_Hour"] = pd.to_datetime(df.Arrival_Time, format="%H:%M").dt.hour  
df["Arr_Min"] = pd.to_datetime(df.Arrival_Time, format="%H:%M").dt.minute  
df["Arrival_Time"] = df['Arr_Hour'] + df['Arr_Min'] / 60  
#df.drop(columns = ['Arr_Hour', 'Arr_Min'], inplace=True)  
df.head()
```

```
# Flight_Duration
```

```
df["FD_Hour"] = df.Flight_Duration.str.split('h').str.get(0)  
df["FD_Min"] = df.Flight_Duration.str.split('h').str.get(1)  
df["FD_Min"] = df["FD_Min"].str.split('m').str.get(0)  
df["FD_Hour"] = df['FD_Hour'].astype('float')  
df["FD_Min"] = df['FD_Min'].astype('float')  
df["Flight_Duration"] = df["FD_Hour"] + df["FD_Min"] / 60  
#df.drop(columns = ["FD_Hour", "FD_Min"], inplace=True)  
df.head()
```

```
# Flight_Prices
```

```
df['Flight_Prices'] = df['Flight_Prices'].str.replace(',', '')  
df['Flight_Prices'] = df['Flight_Prices'].astype('float')  
df
```

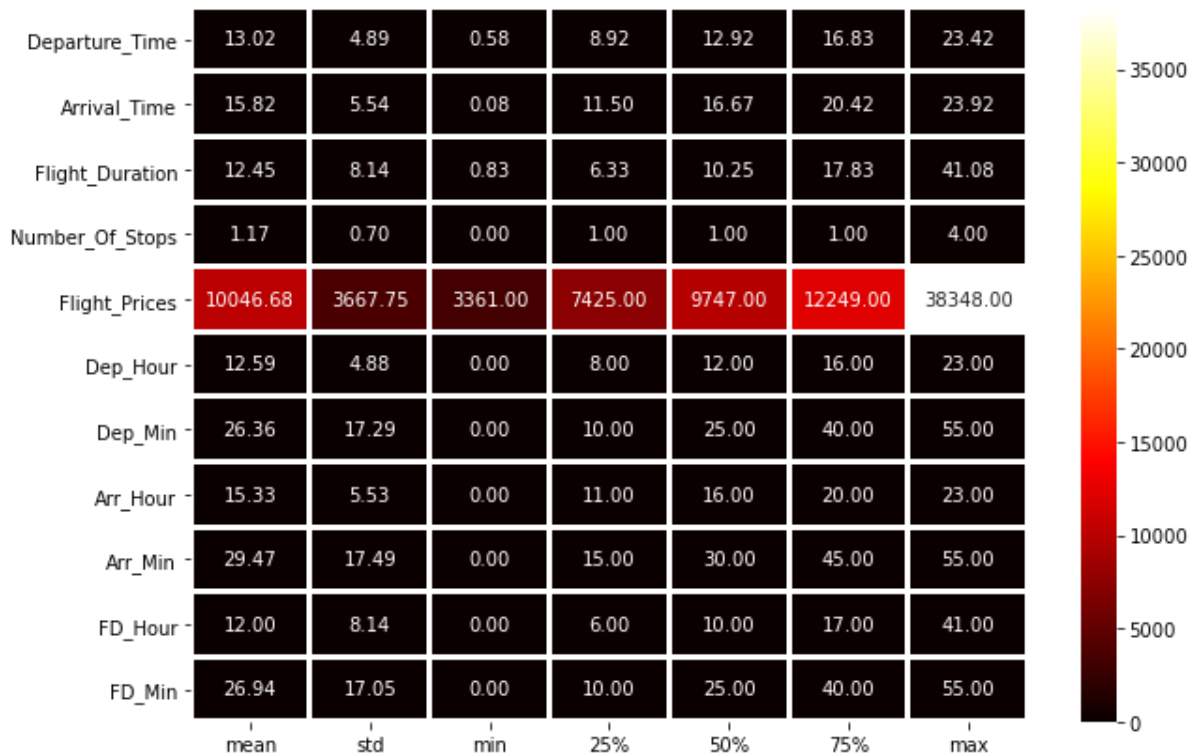
I then used the “describe” method to check the count, mean, standard deviation, minimum, maximum, 25%, 50% and 75% quartile data.

```
df.describe(include="all").T
```

	count	unique	top	freq	mean	std	min	25%	50%	75%	max
Airline_Names	5805	6	IndiGo	1782	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Departure_Time	5805.0	NaN	NaN	NaN	13.024591	4.887243	0.583333	8.916667	12.916667	16.833333	23.416667
Arrival_Time	5805.0	NaN	NaN	NaN	15.823572	5.543619	0.083333	11.5	16.666667	20.416667	23.916667
Flight_Duration	5805.0	NaN	NaN	NaN	12.452728	8.137841	0.833333	6.333333	10.25	17.833333	41.083333
Source_Place	5805	9	Mumbai	806	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Destination_Place	5805	9	Mumbai	805	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Meal_Availability	5805	3	No Meals	4252	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Number_Of_Stops	5805.0	NaN	NaN	NaN	1.172782	0.896018	0.0	1.0	1.0	1.0	4.0
Flight_Prices	5805.0	NaN	NaN	NaN	10046.68062	3667.752804	3361.0	7425.0	9747.0	12249.0	38348.0
Dep_Hour	5805.0	NaN	NaN	NaN	12.585185	4.881136	0.0	8.0	12.0	16.0	23.0
Dep_Min	5805.0	NaN	NaN	NaN	26.364341	17.288619	0.0	10.0	25.0	40.0	55.0
Arr_Hour	5805.0	NaN	NaN	NaN	15.332472	5.526615	0.0	11.0	16.0	20.0	23.0
Arr_Min	5805.0	NaN	NaN	NaN	29.465978	17.493203	0.0	15.0	30.0	45.0	55.0
FD_Hour	5805.0	NaN	NaN	NaN	12.00379	8.137239	0.0	6.0	10.0	17.0	41.0
FD_Min	5805.0	NaN	NaN	NaN	26.936262	17.047469	0.0	10.0	25.0	40.0	55.0

Took a visual on just the numeric part as well and saw just the maximum value for Flight_Prices column at a higher scale.

Statistical Report of Numerical Columns



- Data Inputs- Logic- Output Relationships

The input data were initially all object datatype so had to clean the data by removing unwanted information like “h” and “m” from

Flight_Duration column and ensuring the numeric data are converted accordingly. I then used Ordinal Encoding method to convert all the categorical feature columns to numeric format.

Code:

```
# Ordinal Encoder

oe = OrdinalEncoder()
def ordinal_encode(df, column):
    df[column] = oe.fit_transform(df[column])
    return df

column=["Meal_Availability", "Airline_Names", "Source_Place", "Destination_Place"]
df=ordinal_encode(df, column)
df
```

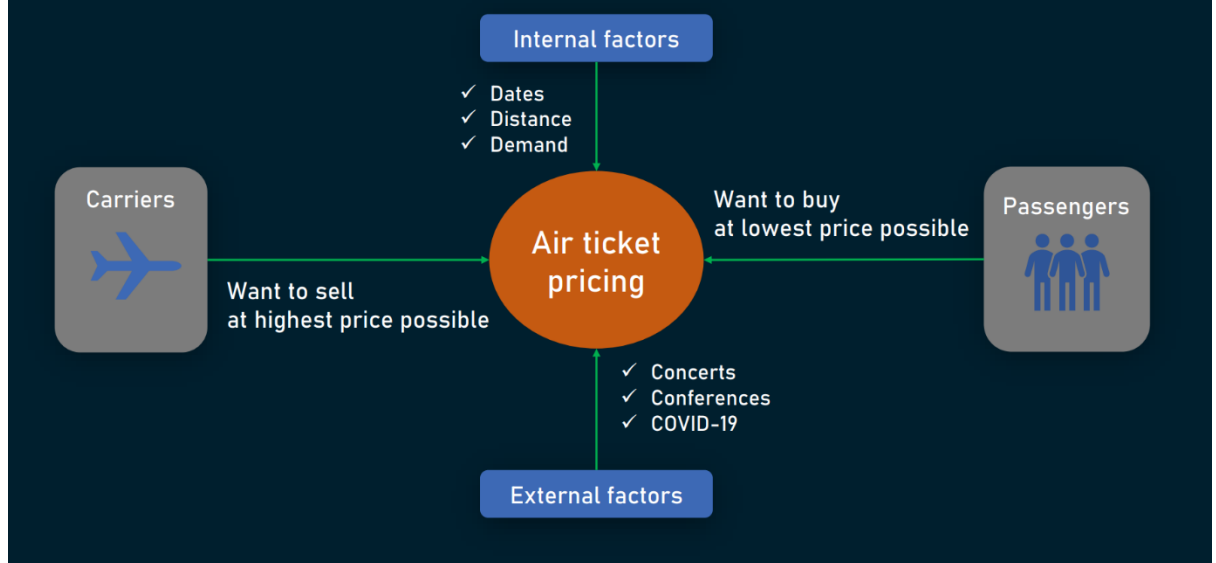
Since we had mostly categorical feature data we did not have to worry about outliers and skewness concerns.

- State the set of assumptions (if any) related to the problem under consideration

Pricing in the airline industry is often compared to a brain game between carriers and passengers where each party pursues the best rates. Carrier's love selling tickets at the highest price possible — while still not losing consumers to competitors. Passengers are crazy about buying flights at the lowest cost available — while not missing the chance to get on board. All this makes flight prices fluctuant and hard to predict. But nothing is impossible for people armed with intellect and algorithms.

There are two main use cases of flight price prediction in the travel industry. OTAs and other travel platforms integrate this feature to attract more visitors looking for the best rates. Airlines employ the technology to forecast rates of competitors and adjust their pricing strategies accordingly.

AIRFARE INFLUENCERS



A passenger-side predictor proposed by an OTA suggests the best time to buy a ticket so that travellers can make informed decisions. Carriers, on their end, try to find out the optimal price they should set to maximize revenue while remaining competitive.

In both cases, the task is quite challenging because numerous internal and external factors influence airfares.

Internal factors include

- purchase and departure dates,
- seasonality,
- holidays,
- the number of available airlines and flights,
- fare class,
- the current market demand, and
- flight distance.

External factors embrace events going on in the arrival or departure cities — like

- concerts,
- sports competitions,
- festivals,

- terrorist attacks,
- natural disasters,
- political gatherings,
- epidemic outbursts, and
- economic activities.

Though it's impossible to cover every external eventuality — say, nothing foreshadowed the 2020 coronavirus pandemic in the middle of 2019 — we still can predict quite a lot, using the right data and advanced machine learning (ML) models.

- **Hardware and Software Requirements and Tools Used**

Hardware technology being used.

RAM : 8 GB

CPU : intel i5 9 Gen

GPU : 4 Graphics NVIDIA GeForce GTX 1050 Ti

Software technology being used.

Programming language : Python

Distribution : Anaconda Navigator

Browser based language shell : Jupyter Notebook

Libraries/Packages specifically being used.

Pandas, NumPy, matplotlib, seaborn, scikit-learn, pandas-profiling, missingno



Model/s Development and Evaluation

- Identification of possible problem-solving approaches (methods)
 1. Clean the dataset from unwanted scraped details.
 2. Rename values with meaningful information.
 3. Encoding the categorical data to get numerical input data.
 4. Compare different models and identify the suitable model.
 5. R2 score is used as the primary evaluation metric.

6. MSE and RMSE are used as secondary metrics.

7. Cross Validation Score was used to ensure there are no overfitting or underfitting models.

- Testing of Identified Approaches (Algorithms)

Libraries and Machine Learning Regression models that were used in this project are shown below.

```
import warnings
warnings.simplefilter("ignore")
warnings.filterwarnings("ignore")
import joblib

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import missingno
import pandas_profiling
from sklearn import metrics
from scipy.stats import zscore
from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import ExtraTreesRegressor
from sklearn.ensemble import GradientBoostingRegressor

from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
```

All the regression machine learning algorithms used are:

- Linear Regression Model
- Ridge Regularization Model
- Lasso Regularization Model
- Support Vector Regression Model
- Decision Tree Regression Model
- Random Forest Regression Model
- K Neighbours Regression Model
- Gradient Boosting Regression Model

- Ada Boost Regression Model
- Extra Trees Regression Model

- Run and Evaluate selected models

I created a Regression Machine Learning Model function incorporating the evaluation metrics so that we can get the required data for all the above models.

Code:

```
# Regression Model Function

def reg(model, X, Y):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25, random_state=638)

    # Training the model
    model.fit(X_train, Y_train)

    # Predicting Y_test
    pred = model.predict(X_test)

    # RMSE - a lower RMSE score is better than a higher one
    rmse = mean_squared_error(Y_test, pred, squared=False)
    print("RMSE Score is:", rmse)

    # R2 score
    r2 = r2_score(Y_test, pred, multioutput='variance_weighted')*100
    print("R2 Score is:", r2)

    # Cross Validation Score
    cv_score = (cross_val_score(model, X, Y, cv=5).mean())*100
    print("Cross Validation Score:", cv_score)

    # Result of r2 score minus cv score
    result = r2 - cv_score
    print("R2 Score - Cross Validation Score is", result)
```

Output:

```
# Linear Regression Model

model=LinearRegression()
reg(model,X, Y)

RMSE Score is: 2741.9142718035005
R2 Score is: 41.30962845018751
Cross Validation Score: 33.96122010005579
R2 Score - Cross Validation Score is 7.34840835013172
```

- Key Metrics for success in solving problem under consideration

RMSE Score:

Root Mean Square Error (RMSE) is the standard deviation of the residuals (prediction errors). Residuals are a measure of how far from the regression line data points are; RMSE is a measure of how spread out these residuals are. In other words, it tells you how concentrated the data is around the line of best fit.

R2 Score:

The R2 score is a very important metric that is used to evaluate the performance of a regression-based machine learning model. It is pronounced as R squared and is also known as the coefficient of determination. It works by measuring the amount of variance in the predictions explained by the dataset.

Cross Validation Score:

Cross-validation is a statistical method used to estimate the skill of machine learning models. It is commonly used in applied machine learning to compare and select a model for a given predictive modelling problem because it is easy to understand, easy to implement, and results in skill estimates that generally have a lower bias than other methods. The k-fold cross validation is a procedure used to estimate the skill of the model on new data. There are common tactics that you can use to select the value of k for your dataset (I have used 5-fold validation in this project). There are commonly used variations on cross-validation such as stratified and repeated that are available in scikit-learn.

Hyper Parameter Tuning:

In machine learning, hyperparameter optimization or tuning is the problem of choosing a set of optimal hyperparameters for a learning algorithm. A hyperparameter is a parameter whose value is used to control the learning process. By contrast, the values of other parameters (typically node weights) are learned.

Code:

```
# Choosing Extra Trees Regressor
```

```
fmod_param = {'n_estimators': [100, 200, 300],  
              'criterion': ['squared_error', 'mse', 'absolute_error', 'mae'],  
              'n_jobs': [-2, -1, 1],  
              'random_state': [42, 251, 340]  
            }
```

```
GSCV = GridSearchCV(ExtraTreesRegressor(), fmod_param, cv=5)  
GSCV.fit(X_train, Y_train)
```

```
GridSearchCV(cv=5, estimator=ExtraTreesRegressor(),  
             param_grid={'criterion': ['squared_error', 'mse', 'absolute_error',  
                                       'mae'],  
                         'n_estimators': [100, 200, 300], 'n_jobs': [-2, -1, 1],  
                         'random_state': [42, 251, 340]})
```

Final model score after plugging in the best parameter values:

```
Final_Model = ExtraTreesRegressor(criterion='mae', n_estimators=300, n_jobs=-2, random_state=251)  
Model_Training = Final_Model.fit(X_train, Y_train)  
fmod_pred = Final_Model.predict(X_test)  
fmod_r2 = r2_score(Y_test, fmod_pred, multioutput='variance_weighted')*100  
print("R2 score for the Best Model is:", fmod_r2)
```

R2 score for the Best Model is: 72.4377190016525

- Visualizations

I used the pandas profiling feature to generate an initial detailed report on my dataframe values. It gives us various information on the rendered dataset like the correlations, missing values, duplicate rows, variable types, memory size etc. This assists us in further detailed visualization separating each part one by one comparing and research for the impacts on the prediction of our target label from all the available feature columns.

Code:

```
pandas_profiling.ProfileReport(df)
```

Output:

Summarize dataset: 100%  28/28 [00:25<00:00, 1.38s/it, Completed]

Generate report structure: 100%  1/1 [00:06<00:00, 6.65s/it]

Render HTML: 100%  1/1 [00:04<00:00, 4.28s/it]

Pandas Profiling Report

[Overview](#)[Variables](#)[Interactions](#)[Correlations](#)[Missing values](#)[Sample](#)[Duplicate rows](#)

Overview

[Overview](#)[Warnings](#) 38[Reproduction](#)

Dataset statistics


Number of variables	15
Number of observations	5805
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	232
Duplicate rows (%)	4.0%
Total size in memory	680.4 KiB
Average record size in memory	120.0 B

Variable types

Categorical	5
Numeric	10

I generated count plots, bar plots, pair plots, heatmap and others to visualise the datapoint present in our column records.

Count Plot for Airline_Names column



Observation on Fig 1:

There are 73.25% flights that mostly serve no meals in their domestic journey since they are of short distances and duration

We can see that 17.47% flights serve free meals which are probably for tickets that include those prices and meal services

Finally, there are 9.29% flights which offer the eCash meals option that can be redeemed to purchase food during long journey flights mostly with multiple stops

Observation on Fig 2:

Highest number of airlines preferred by people are Indigo covering 30.7% of the total record

We can see that Air India is quite close to the first one and a close competitor standing at the second position holding 28.22% of the total record

At third place we have Vistara airlines that covers 26.08% of total record in our airline data

Airlines Go First, Air Asia and SpiceJet are the least used by people covering 8.48%, 3.67% and 2.86% respectively

Observation on Fig 3:

The departure area or source place highly used or people majorly flying from the city is "Mumbai" covering 13.88% record in the column

We see that "Bangalore" is a close second wherein it covers 13.01% records in the column

Other two famous locations where people chose to fly from are "New Delhi" and "Kolkata" covering 12.89% and 12.82% respectively

The least travel from location is "Jaipur" where I believe not many people chose to travel from unless they are returning from a vacation there and hence covers 7.36% of record in the column

Observation on Fig 4:

Just as in case of departure area even in terms of arrival area or destination place people prefer to fly towards the city "Mumbai" covering 13.87% of record

Again, in a similar fashion "Bangalore" city is a close second destination that people like to fly towards covering 13.54% record in the column

Surprisingly we have "Hyderabad" city taking the third place for destination that people prefer landing covering 13.07% of record

Finally, similar to the departure location the least travel to area is "Jaipur" and it covers 5.24% record in the column

Observation on Fig 5:

People generally buy flight tickets that have 1 stop layover covering 63.67% rows in the column

Next in line are 2 stop layovers which cover 19.88% rows

Here we can see that the 0 stop flights availability is around 11.94% of the total record


In domestic flight we rarely see 3 or 4 stop layovers and therefore they cover 4.2% and 0.31% of total rows in the column respectively

Code:

```
x = "Source_Place"
plt.figure(figsize=(18,10))
sns.countplot(x = x, hue = "Airline_Names", data = df)
plt.title(f"Countplot for {x} column\n", fontsize = 20)
plt.show()

x = "Destination_Place"
plt.figure(figsize=(18,10))
sns.countplot(x = x, hue = "Airline_Names", data = df)
plt.title(f"Countplot for {x} column\n", fontsize = 20)
plt.show()
```

Output:



Observation:

Checking out the Source place details for each and every airline we can see that Mumbai city has the highest number of departure flights for Air India airlines

Go First, Indigo and Air India are the airlines that are used in almost all the cities to depart while the other airlines do not cover some or the other city

Looking at the Destination place details for each and every airline we can see that Hyderabad city has the highest number of arrival flights for Air India airlines

Once again, we can observe that Go First, Indigo and Air India are the leading airlines that are used in almost all cities to arrive while the other airlines miss out on some or the other regions

Overall, I can notice that Air India and Indigo flights do quite well and can be used for arrival and departure to and from any location in India

Code:

```

y = 'Airline_Names'

x = 'Departure_Time'
plt.figure(figsize=[15,7])
sns.barplot(x,y,data=df,orient='h')
plt.title(f"Barplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()

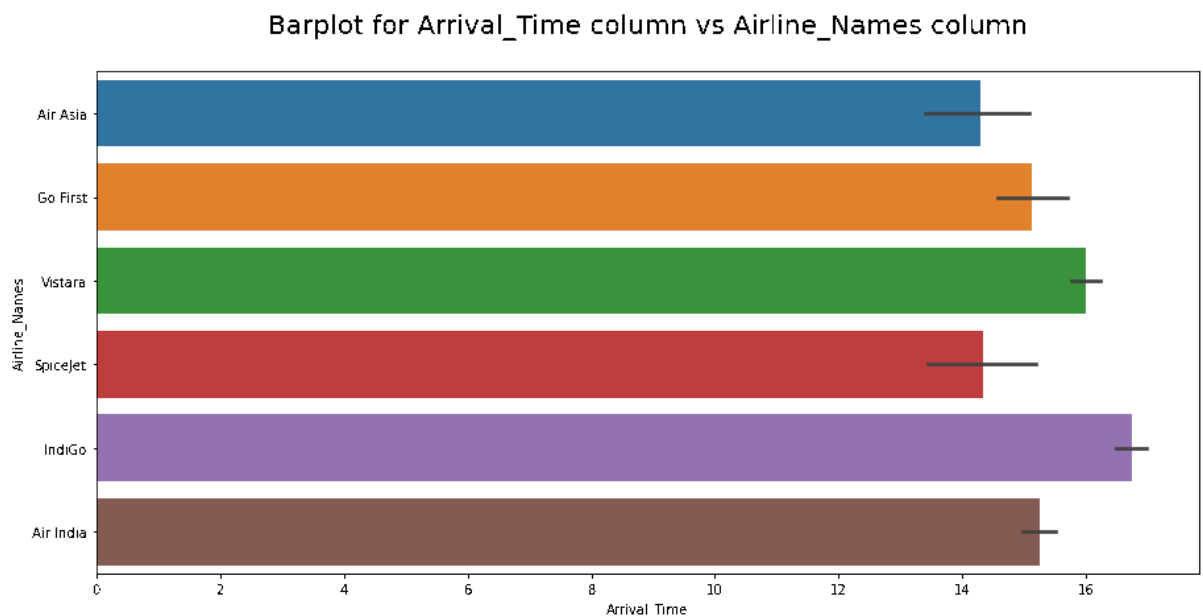
x = 'Arrival_Time'
plt.figure(figsize=[15,7])
sns.barplot(x,y,data=df,orient='h')
plt.title(f"Barplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()

x = 'Flight_Duration'
plt.figure(figsize=[15,7])
sns.barplot(x,y,data=df,orient='h')
plt.title(f"Barplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()

x = 'Flight_Prices'
plt.figure(figsize=[15,7])
sns.barplot(x,y,data=df,orient='h')
plt.title(f"Barplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()

```

Output:



Observation:

When we observe the barplot for Departure time vs Airline we can see that Air Asia has the highest departure time while IndiGo has the lowest departure time

Considering the barplot for Arrival time vs Airline we can see that IndiGo has the highest arrival time while Air Asia and SpiceJet have the lowest arrival time

Taking a look at the barplot for Flight duration vs Airline we observe that Ai India has the highest flight duration while Indigo has the lowest flight duration collectively

Comparing the bar plots for Flight prices vs Airline we can clearly see that Vistara and Air India have very high flight prices while the other airlines IndiGo, SpiceJet, Go First and Air Asia lie in a similar price bracket where Air Asia has the lowest fare

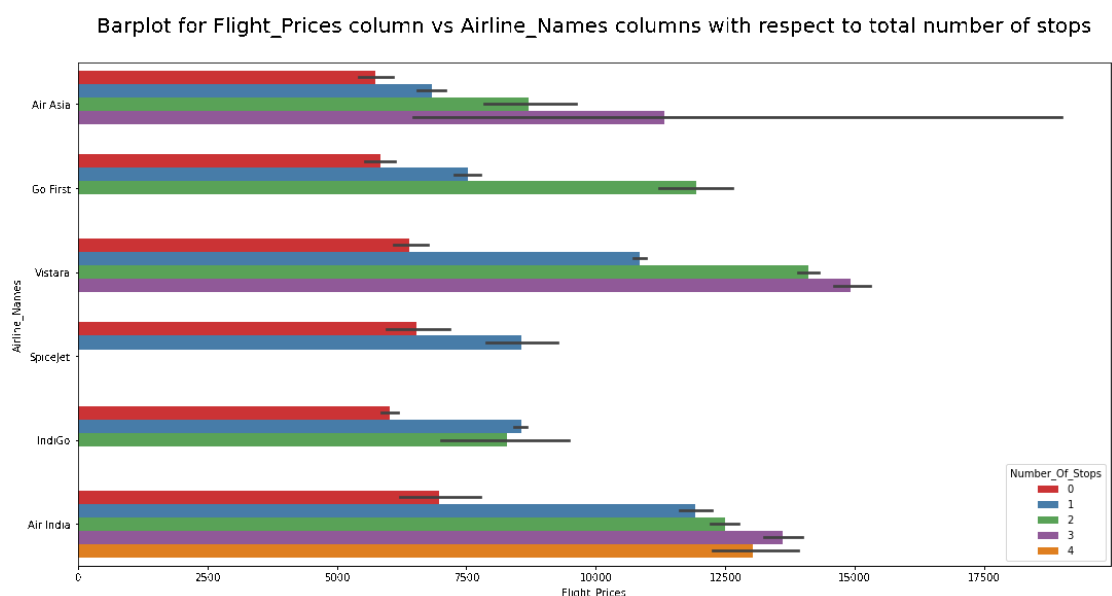
Code:

```
x = "Flight_Prices"
y = "Airline_Names"

plt.figure(figsize=(18,9))
sns.barplot(x=df[x], y=df[y], hue=df['Meal_Availability'], palette="Set1")
plt.title(f"Barplot for {x} column vs {y} columns with respect to total number of stops\n", fontsize = 20)
plt.show()

plt.figure(figsize=(18,9))
sns.barplot(x=df[x], y=df[y], hue=df['Number_Of_Stops'], palette="Set1")
plt.title(f"Barplot for {x} column vs {y} columns with respect to total number of stops\n", fontsize = 20)
plt.show()
```

Output:



Observation:

Vistara and Air India are the only airlines that have free meal service on their flights but also have a very high no meal record

Also, the eCash meal options do not work for Vistara and Air India flights and they have very high flight prices compared to their competitors

If we see a trend the flights with 0 stops or rather direct flights for every airline is cheaper when compared to 1 or more layovers

And flights with 2 and 3 stops have a considerably high price and number of flights available in those records are high too

Only Air India provides flights with 4 stop layovers even in a domestic environment

SpiceJet on the other hand only has flights available with 0 and 1 stop layovers

Code:

```
y = "Flight_Prices"

x = "Airline_Names"
plt.figure(figsize = (15,8))
sns.barplot(data = df, y = y, x = x)
plt.title("Prices according to different Airlines\n", fontsize = 20)
plt.show()

x = "Source_Place"
plt.figure(figsize = (15,8))
sns.barplot(data = df, y = y, x = x)
plt.title("Prices according to different Source places\n", fontsize = 20)
plt.show()

x = "Destination_Place"
plt.figure(figsize = (15,8))
sns.barplot(data = df, y = y, x = x)
plt.title("Prices according to different Destination places\n", fontsize = 20)
plt.show()

x = "Number_of_Stops"
plt.figure(figsize = (8,8))
sns.barplot(data = df, y = y, x = x)
plt.title("Prices according to different Number of layover stops\n", fontsize = 20)
plt.show()

x = "Meal_Availability"
plt.figure(figsize = (8,8))
sns.barplot(data = df, y = y, x = x)
plt.title("Prices according to different Meal options\n", fontsize = 20)
plt.show()
```

Output:



Observation:

Airfares in Vistara and Air India are pretty high when compared to other airlines as they provide free meal service which probably is just meal cost included in the tickets

Flight prices when departing from cities like Kolkata and Chennai have higher price range but the others are around the similar range a bit lesser in pricing but not providing a huge difference as such

Similarly, prices when arriving from cities Kolkata and Chennai have high price range however the highest recorded is for Lucknow city and the rest fall in similar price bracket again not making any huge difference in savings

When we consider the layovers for pricing situation then obviously direct flights are cheaper when compared to flights that have 1 or more stops

As per the data collected, we see that free meal service flight tickets cost the highest as compared to no meal service flight tickets and the lowest price observed are for eCash meal service wherein

the frequent flyers can use their flight points as discounts on the ticket prices

Code:

```
y = "Flight_Prices"


x = "Departure_Time"
plt.figure(figsize = (15,8))
sns.scatterplot(x=x,y=y,data=df)
plt.title(f"Scatterplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()

x = "Arrival_Time"
plt.figure(figsize = (15,8))
sns.scatterplot(x=x,y=y,data=df)
plt.title(f"Scatterplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()

x = "Flight_Duration"
plt.figure(figsize = (15,8))
sns.scatterplot(x=x,y=y,data=df)
plt.title(f"Scatterplot for {x} column vs {y} column\n", fontsize = 20)
plt.show()
```

Output:

Scatterplot for Arrival_Time column vs Flight_Prices column



Observation:

In the above scatter plot, we see the comparison where flight prices are the highest during peak hours as compared to late hours or specifically graveyard timings for departure


We can see a similar trend in flight price scatter plot details when it comes to arrival timings where we see a decrease in prices between 1:00 hrs and 6:00 hrs

Finally, as the overall flight duration increases the flight prices increase too and that makes direct flight depart and arrive in a relatively shorter period of time

Code:

```
print("***** Pair Plot with Meal Type Legend *****")
sns.pairplot(df, hue='Meal_Availability', diag_kind="kde", kind="scatter", palette="Set1", height=3.5)
plt.show()
```

Output:




Code:

```
plt.style.use('seaborn-bright')

df.hist(figsize=(20,20))
plt.show()
```


Output:



Code:

```
upper_triangle = np.triu(df.corr())
plt.figure(figsize=(15,10))
sns.heatmap(df.corr(), vmin=-1, vmax=1, annot=True, square=True, fmt='0.3f',
            annot_kws={'size':10}, cmap="magma", mask=upper_triangle)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
plt.show()
```


Output:



Code:

```
df_corr = df.corr()
plt.figure(figsize=(14,7))
df_corr['Flight_Prices'].sort_values(ascending=False).drop('Flight_Prices').plot.bar()
plt.title("Correlation of Feature columns vs Label column", fontsize=20)
plt.xlabel("\nFeatures List", fontsize=14)
plt.ylabel("Correlation Value", fontsize=14)
plt.show()
```

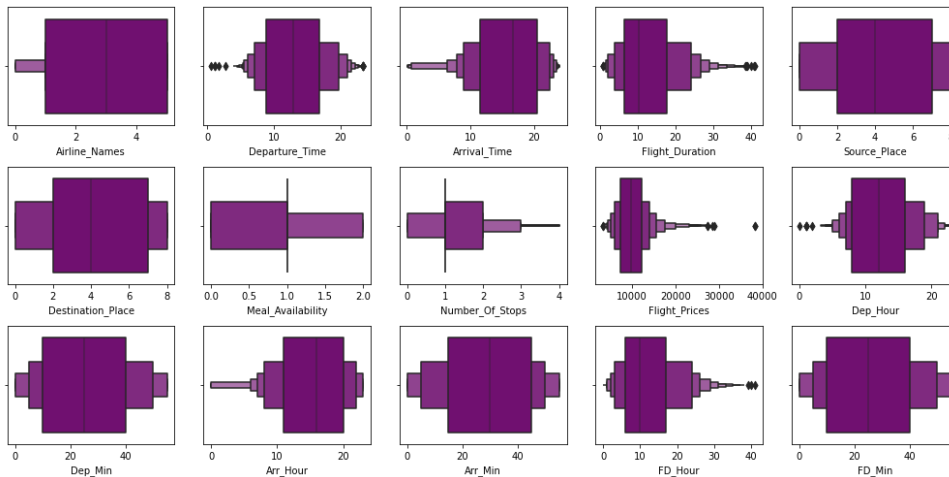
Output:



Outliers Code:

```
plt.figure(figsize=(14,7))
outl_df = df.columns.values
for i in range(0, len(outl_df)):
    plt.subplot(3, 5, i+1)
    ax = sns.boxenplot(df[outl_df[i]], color='purple')
    plt.tight_layout()
```

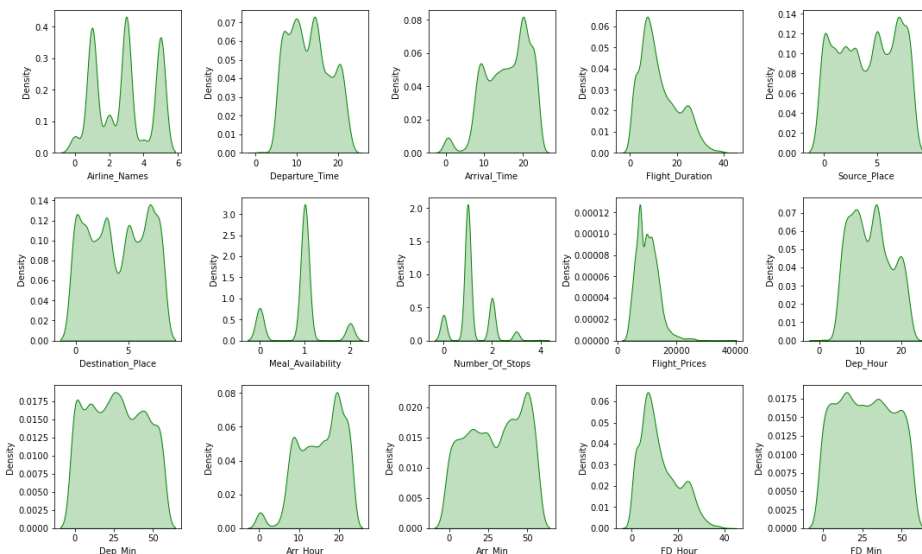
Output:



Skewness Code:

```
fig, ax = plt.subplots(ncols=5, nrows=3, figsize=(15,9))
index = 0
ax = ax.flatten()
for col, value in df.items():
    sns.distplot(value, ax=ax[index], hist=False, color="g", kde_kws={"shade": True})
    index += 1
plt.tight_layout(pad=0.8, w_pad=0.8, h_pad=2.0)
plt.show()
```


Output:



Importance Graph Code:

```
rf=RandomForestRegressor()  
rf.fit(X_train, Y_train)  
importances = pd.DataFrame({'Features':X.columns, 'Importance':np.round(rf.feature_importances_,3)})  
importances = importances.sort_values('Importance', ascending=False).set_index('Features')  
plt.rcParams["figure.figsize"] = (15,8)  
importances.plot.bar(color='teal')  
importances
```

Output:



- Interpretation of the Results

From the above EDA we can easily understand the relationship between features and we can even see which things are affecting the price of flights. These methods take financial, marketing, and various social factors into account to predict flight prices.

Nowadays, the number of people using flights has increased significantly. It is difficult for airlines to maintain prices since prices change dynamically due to different conditions. That's why we have tried to use machine learning to solve this problem. This can help airlines by predicting what prices they can maintain. It can also help customers to predict future flight prices and plan their journey accordingly.

CONCLUSION

- Key Findings and Conclusions of the Study

In this project we have scraped the flight data from airline webpages. Then the comma separated value file is loaded into a data frame. Luckily, we don't have any missing values in our data set. Looking at the data set we understand that there are some features needs to be processed like converting the data types and get the actual value from the string entries from the time related columns. After the data is been processed, I have done some EDA to understand the relation among features and the target variable. Features like flight duration, number of stops during the journey and the availability of meals are playing major role in predicting the prices of the flights. As we have seen, the prediction is showing a similar relationship with the actual price from the scrapped data set. This means the model predicted correctly and it could help airlines by predicting what prices they can maintain. It could also help customers to predict future flight prices and plan the journey accordingly because it is difficult for airlines to maintain prices since it changes dynamically due to different conditions. Hence by using Machine Learning techniques we can solve this problem. The above research will help our client to study the latest flight price market and with the help of the model built he can easily predict the price ranges of the flight, and also will helps him to understand Based on what factors the fight price is decided.

- Learning Outcomes of the Study in respect of Data Science

Visualization part helped me to understand the data as it provides graphical representation of huge data. It assisted me to understand the feature importance, outliers/skewness detection and to compare the independent-dependent features. Data cleaning is the most important part of model building and therefore before model building, I made sure the data is cleaned. I have generated multiple

regression machine learning models to get the best model wherein I found Extra Trees Regressor Model being the best based on the metrics I have used.

